

TIP1: Linux Dev Tools/Tips for

C/C++ Debugging/Tracing/Profiling

Agenda

- Preface
- Concepts
- Tools for C/C++
 - Debugging
 - Tracing
 - Profiling
- References
- Postscript

Preface

- What does our world look like?

"There is no remembrance of former things; neither shall there be any remembrance of things that are to come with those that shall come after."

-- Ecclesiastes 1:11

We want/need/have to change this...

TIP = Technology Inheritance Program

Concepts

- *Debugging* - Find the cause of unexpected program behavior, and fix it.
- *Profiling* - Analyze program runtime behavior, provide statistical conclusions on key measurements (speed/resource/...).
- *Tracing* - Temporally record program runtime behavior, provide data for further debugging/profiling.

All debugging/profiling/tracing tools depend on some kind of **instrumentation** mechanism, either statical or dynamical.

Debugging tools for C/C++

Debugging Tools Implementation

- Breakpoint support
 - Hardware breakpoint
 - **DR0~7** regs on Intel CPU
 - Software breakpoint
 - **INT3** instruction on x86/x86_64
 - raise **SIGTRAP** signal for portable breakpoint
 - Virtual Machine Interpreter
 - Interpret instructions instead of execute it directly
- Linux user-space debug infrastructure
 - `ptrace` `syscall`

Debugging Tools

- **gdb** - General-purpose debugger
 - ptrace-based
 - Both hw/sw breakpoints supported
 - **Reverse executing** feature in 7.x version
 - Save reg/mem op before each instr executed, heavy but very handy
 - Usecases:
 - Standalone debug
 - `gdb --args <exec> <arg1> <...>`
 - Analyze core
 - `gdb <exec> <core>`
 - Attach to existing process
 - `gdb <exec> <pid>`
 - Many resources, search and learn:)

Debugging Tools

- **Valgrind** family

- valgrind is an instruction interpreter/vm framework
- **Impossible** to attach to a running process :(
- Useful plugin:

- **memcheck**

- Memory error detector

- **massif**

- Heap usage profiler

- **helgrind**

- Thread error detector

- **DRD**

- (another) Thread error detector

- **ptrcheck(SGCheck)**

- Stack/global array overrun detector

Debugging Tools

- memcheck usecases:
 - Check memory error for all process in hierarchy:
 - `valgrind --tool=memcheck --leak-check=full --leak-resolution=high --track-origins=yes --trace-children=yes --log-file=./result.log <exec>`
 - See flags specified to memcheck plugin:
 - `valgrind --tool=memcheck --help`

Debugging Tools

- massif usecases:
 - Stats heap and stack usage during a program's life:
 - `valgrind --tool=massif --stacks=yes <exec>`
 - `ms_print massif.*`
 - In the output of `ms_print`:
 - `:` means normal snapshot
 - `@` means detail snapshot
 - `#` means peak snapshot in all

Debugging Tools

- helgrind usecase:
 - Check POSIX thread API misuse/inconsistent lock order/data races:
 - `valgrind --tool=helgrind <exec>`
- DRD usecase:
 - Check POSIX thread API misuse/data races/lock contention, and tracing all mutex activities:
 - `valgrind --tool=drd --trace-mutex=yes <exec>`
- ptrcheck usecase:
 - Check stack/global array overrun:
 - `valgrind --tool=exp-ptrcheck <exec>`

Debugging Tools

- Intel Inspect XE (Commercial)
 - Cross-platform proprietary debugging tools
 - Both GUI/CLI usage supported
 - Memory/thread error detector
 - Free for non-commercial use
 - Included in Intel Parallel Studio suite, standalone download available
 - **Catch up very slow on new hardwares (e.g. i7...)**
 - **Works not well on Linux platform, other platform not tested...**

Debugging Guideline

- Generally speaking, all programs should pass Valgrind **memcheck/ptrcheck** checking, to eliminate most of the memory errors.
- Multithread programs should pass Valgrind **helgrind/drd** checking, to eliminate common racing errors.
- Valgrind **massif** can be used to track down the origin of unexpected heap allocation.
- **gdb** can be used to manually track down logical bugs in the code.
- Multiprocess/thread programs don't fit gdb well, most of the time tracing the program is much easier/faster to find the source of a bug than manually gdb debugging.

Profiling tools for C/C++

Profiling Tools Implementation

- Event based profiling
 - Add hook for specified event, count event occurring times
- Sampling based profiling
 - Make a repeating trigger for sampling
 - Record instruction counter and call stack when trigger'd
 - Generate statistically result based on record data

NOTE: General profiling tools can **NOT** reveal **sleeping** (interruptible blocking, lock wait, etc.) or **I/O blocking** (non-interruptible blocking) costs! But these are usually the **main throttle** to the intuitive runtime performance.

Profiling Tools (event based)

- **gcov**

- A coverage testing tool, but can also be used as a line-count profiling tool (user-space only)
- Need statistically instrument target program, compiling with one of the following gcc flags:
 - `--coverage`
 - `-fprofile-arcs -ftest-coverage`
- When program exits normally, *.gcda/gcno file will be generated
- Usecase:
 - `gcc --coverage x.c -ox`
 - `gcov x.c # gen x.c.gcov`
 - `less x.c.gcov`

Profiling Tools (event based)

Behind the scene of **gcov**:

- **-ftest-coverage** makes compiler generating *.gcnf files, which contains infos to reconstruct basic block graph and assign source codes to blocks (used by gcov).
- **-fprofile-arcs** makes compiler injecting codes adding counters associated with each source code line, and codes that dump out *.gcda files when the program exits.
- See:
 - `gcc -S x.c -o x1.s`
 - `gcc -S --coverage x.c -o x2.s`
 - `vimdiff *.s`

Profiling Tools (event based)

- **lcov**

- Graphical gcov front-end
- Generate beautiful coverage report in HTML format
- Usecase:

- Assuming the source is placed in app/x.c

- `cd app`

- `gcc --coverage x.c -ox`

- `./x`

- `lcov -d . -c -o x.info`

- `genhtml -o report x.info`

- See `app/report/index.html` for report

Profiling Tools (event based)

- **valgrind (callgrind)**

- Instruction level profiler, with cool GUI frontend

- kcachegrind**

- Cache/branch prediction profiling and annotated source supported

- Add -g compiler flag if annotated source is wanted

- Usecase:

- `gcc -g x.c -o x`

- `valgrind --tool=callgrind --dump-instr=yes --cache-sim=yes --branch-sim=yes ./x`

- `kcachegrind callgrind.*`

Profiling Tools (sampling based)

- **gprof**

- Timer based IP sampling + call event count
 - Use `setitimer(ITIMER_PROF, ...)` on Linux
 - Sampling frequency depends on kernel's HZ setting
- Flat report, call graph report and annotated source supported
- Compiling & Linking with flag `-pg`
 - Add `-g` if annotated source is wanted
- Usecase:
 - `gcc -pg -g x.c -o x`
 - `./x # gmon.out gen'd`
 - `gprof ./x # see flat/call graph report`
 - `gprof -A ./x # see annotated source`

Profiling Tools (sampling based)

Behind the scene of **gprof**:

- gprof is supposed to use `profil()` syscall for IP sampling, but that syscall is not implemented by Linux kernel, so it falls back to mimic the syscall with `setitimer()`.
- **-pg** makes compiler injecting codes calling `mcount()` at the entry of each function, which collects call-stack info.
 - `gcc -S x.c -ox1.s`
 - `gcc -S -pg x.c -ox2.s`
 - `vimdiff *.s`
- This options also makes linker linking with `gcrt*.o` instead of normal `crt*.o`, which provides startup routine to init sampling timers and resources.
 - `gcc -v x.c | grep crt`
 - `gcc -v -pg x.c | grep crt`

Profiling Tools (sampling based)

- **google-perftools (CPU profiler)**

- Timer based call-stack sampling
 - Use `setitimer(ITIMER_PROF, ...)` on Linux
 - Set sampling frequency through env var `PROFILEFREQUENCY`
- Linked-in usage (NOTE: profiler symbols must be referenced in your code, otherwise the dependency of profiler shared library will be eliminated!)
 - `gcc -g x.c -ox -lprofiler`
 - `CPUPROFILE=/tmp/xxx ./x`
- Preload usage:
 - `LD_PRELOAD=/usr/local/lib/libprofiler.so`
`CPUPROFILE=/tmp/xxx ./x`
- Show report: `pprof --text ./x /tmp/xxx`

Profiling Tools (sampling based)

- **oprofile**

- Support timer/interrupt/PMC/tracepoint based sampling
 - PMC = PerformMance Counter
- Capable of doing system-wide profiling
- Deprecated in prefer of **perf** on kernel > 2.6.26(?)
- Usecase:
 - `sudo opcontrol --init # load oprofile module`
 - `sudo opcontrol -s`
 - `./x`
 - `sudo opcontrol -h`
 - `sudo opreport # show report`
 - `sudo opannotate -s # show annotated src`

Profiling Tools (sampling based)

- **perf**

- Available on kernel $\geq 2.6.26(?)$
- PMC frontend released along with kernel itself
- Support PMC/tracepoint based sampling
- Capable of doing system-wide profiling, sampling events trace can also be output
- Usecase:
 - `sudo perf record -a -g -- ./x`
 - `sudo perf report # show prof report`
 - `sudo perf annotate # show annotated src`

Profiling Tools (sampling based)

- **Intel VTune Amplifier XE (Commercial)**

- PMC/timer based sampling, support GUI/CLI
- System-wide profiling supported, has locks & waits analysis
- Use **Pin** for instrumentation
- CLI works well on Linux, **GUI not stable**
 - `amplxe-cl -collect hotspots ./x`
 - `amplxe-cl -report hotspots -r rxxxxhs`

- **AMD CodeAnalyst (Commercial)**

- **oprofile** based, GUI only
- System-wide profiling supported
- Provide much more events on AMD CPUs
- **Works not well on Linux**

Profiling Guideline

- Determine target program performance throttle before actual profiling (time helps)
 - sys time + user time \sim wall clock time
 - sys time \gg user time: reduce syscalls / user-kernel space profiling
 - user time \gg sys time: user space profiling
 - sys time + user time \ll wall clock time
 - Don't use general profiling tool, consider user space tracing
- Analysis profiling result hierarchically, starting from outter scope first, don't dive into details too early.
- Spot performance throttle one by one. First deal with the biggest known throttle, then profiling again and find the next throttle.

Tracing tools for C/C++

Tracing Tools Implementation

- Decouple event recording and exporting: ring buffer
- User-space tracing
 - Intrusive
 - Call tracing API manually, need recompiling code
 - Non-intrusive
 - ptrace syscall
 - GNU dynamic linker LD_AUDIT
 - utrace-patched kernel
- Kernel-space tracing
 - Dynamical mechanism
 - kprobes / jprobes / kretprobes: trap/short-jmp instr
 - Statical mechanism
 - tracepoints: manually inserted conditional jump
 - ftrace (kernel \geq 2.6.26): gcc mcount utilization

Tracing Tools (ptrace based)

- **strace**

- Trace user program's syscalls
- Support existing process tracing
 - Watch out ptrace protection patch! (for nonroot) `/proc/sys/kernel/yama/ptrace_scope`
- Works well with multithread programs
- Usecase:
 - `strace -f -i -tt -T -v -s 1024 -C -o trace.out ./x`
 - See `man strace` for detail description

Tracing Tools (ptrace based)

- **ltrace**

- Trace user program's dynamic library calls
- Can also trace syscalls, but can't parse their args as strace did
- Neither library->library nor dlopen'd library call trace supported
- **Can NOT work with multithread programs**
- Usecase:
 - `ltrace -C -f -i -n4 -s1024 -S -tt -T ./x`
 - See man ltrace for detail description

Tracing Tools (ptrace based)

- Ptrace-based tracing shortcoming:
 - **Heavy overhead**, at least **2 ctx sw + 2 syscall** plus **signal transit** overheads per tracepoint, very slow on large tracepoint set;
 - `init(1)` can not be traced;
 - Processes can not be ptraced by multiple tracers;
 - Ptrace **affects the semantics of traced processes**:
 - Original parent will not be notified when its child was ptraced and stopped (see notes in `man 2 ptrace`)
 - The overhead of ptrace will lower the num of concurrent running threads. Race conditions sensitive to timings may disappear due to this, resulting a Heisenberg problem.

Tracing Tools (LD_AUDIT based)

- **latrace**

- Trace user program's dynamic library calls
- Can NOT trace existing process
- Use callback function running in target process instead of ptrace signals, much lower overhead
- Works well with multithread programs
- Usecase:
 - `latrace -SAD -o trace.out ./x`
 - See man latrace for detail description

Tracing Tools (ftrace based)

- **trace-cmd**

- Available on kernel $\geq 2.6.26$
- CLI frontend for ftrace framework
- System-wide kernel tracer, no user space event available (except for events like context switching, scheduling etc., but no call-site info)
- Usecase:
 - `sudo trace-cmd record -e all -p function_graph -F ./x`
 - `trace-cmd report`

Tracing Tools (ftrace based)

- **kernelshark**

- GUI viewer for trace-cmd result

- Usecase:

- `sudo trace-cmd record -e all -p function_graph -F ./x`

- `kernelshark`

Tracing Tools (customized)

- **SystemTap**

- Linux community's reply to Solaris DTrace
- Scriptable framework to utilize kprobes/tracepoints
- User space tracing needs utrace-patched kernel, Redhat distros (RHEL/CentOS/Fedora) all comes with such kernels
- Usecase:
 - `stap -e 'probe syscall.* {println(thread_indent(4), "->", probefunc())} probe syscall.*.return {println(thread_indent(-4), "<-", probefunc())}' -c ./x`

Tracing Tools (customized)

- **LTTng 2.0**

- Rewrite of LTTng 0.9.x, no need to patch kernel anymore, lighter weight compare to SystemTap
- User space tracing is done by inserting statical tracepoint into user program (not compatible with SystemTap/DTrace probes yet...)
- Usage:
 - `sudo lttng create sess1`
 - `sudo lttng enable-event -a -k`
 - `sudo lttng enable-event -a -u`
 - `sudo lttng start`
 - `./x`
 - `sudo lttng stop`
 - `babeltrace ~/lttng/sess1*`

Tracing Tools (customized)

- **DTrace**

- Origins from Sun Solaris, adopted by MacOS/FreeBSD/Oracle Unbreakable Linux
- Scriptable framework, light weight tracing overhead
- Capable of kernel and user space joint tracing (user space tracing needs inserting statical tracepoints)
- Handy tracing multiple languages / apps:
 - Java(Sun) / PHP(Zend) / Javascript(Firefox) / CPython / CRuby / MySQL / PostgreSQL / Erlang (DTrace fork)
- Usecase: see <http://dtracehol.com/>

Tracing Guideline

- **Be warned!** Tracing needs involved efforts and solid background on Linux kernel. Learn more and deeper about how the system working first!
- Use SystemTap for kernel/user space tracing on Redhat family distros (RHEL/CentOS/Fedora) or utrace-patched kernels
- Use DTrace for kernel/user space tracing on MacOS/FreeBSD
- User space only tracing can be partially done by strace/ltrace
- LTTng 2.0 can do kernel/user space tracing if you can insert statical tracepoints in your code, and it does not need patching your kernel

Practical Tips

Other useful technics

- gcc -finstrument-functions
 - <https://github.com/agentzh/dodo/tree/master/utils/dodo-hook>
- LD_PRELOAD crash signal handler
 - https://github.com/chaoslawful/phoenix-nginx-module/tree/master/misc/dbg_jit
- Add signal handler to normally output gprof/gcov result for a interrupted program

See examples at <https://github.com/chaoslawful/TIP>

References

- Overview

- [Linux Instrumentation](#)
- <http://lwn.net/Kernel/Index/>

- Tracing

- [玩转 utrace](#)
- [utrace documentation file](#)
- [Introducing utrace](#)
- [Playing with ptrace, Part I](#)
- [Playing with ptrace, Part II](#)
- [SystemTap/DTrace/LTTng/perf Comparison](#)
- [ftrace 简介](#)
- [Solaris Dynamic Tracing Guide](#)
- [DTrace for Linux](#)
- [Observing and Optimizing your Application with DTrace](#)

References

- Tracing

- [SystemTap Beginner's Guide](#)
- [SystemTap Language Reference](#)
- [SystemTap Tapset Reference](#)
- [LTTng recommended bundles](#)
- [LTTng Ubuntu daily PPA](#)
- [An introduction to KProbes](#)
- [使用KProbes调试内核](#)
- [Tracing: no shortage of options](#)
- [Uprobes: 11th time is the charm?](#)
- [Ptrace, Utrace, Uprobes: Lightweight, Dynamic Tracing for User Apps](#)
- [LTTng Tracing Book](#)

References

- Profiling & Debugging
 - [google-perftools Profiling heap usage](#)
 - [google-perftools CPU Profiler](#)
 - [Valgrind User Manual](#)
 - [OProfile Manual](#)
 - [Debugging with GDB](#)
 - [GDB Internals Manual](#)
 - [Implementation of GProf](#)
 - [Gcov Data Files](#)

Postscript

*"The important thing is **not to stop** questioning; **never lose** a holy curiosity."*

-- Albert Einstein